

AuthorsLab × High Line Publishing

A proposed infrastructure partnership for a new publishing operation

WORKING DRAFT V0.5 · INTERNAL — NOT FOR SEND · this is the voice-pass and countersign text; one section awaits its owner and is marked. Final version follows countersign, adversarial read and verification.

High Line is building a new publishing operation — and the systems, workflows and infrastructure around its author list, at the same time as the list itself. This proposal positions AuthorsLab not as a software purchase but as an **infrastructure partner**: a connected environment for authors, editorial teams and publishing management, adopted progressively, with minimal operational or financial friction, that develops alongside the publishing operation rather than requiring High Line to build every component internally.

1 · The problem, in your words

“We have taken on a lot of new authors and we need to get them to market as soon as possible.”

That is a throughput problem, and throughput problems have a particular shape: the cost of a slow book is not the editing bill — it is the season it misses. The question that needs answering every week, for every title, is not “*how is it going?*” but “**which book is going to slip?**” Everything in this proposal exists to make you more certain about when your books reach market.

2 · The proposition

One environment around the life of a book: manuscript development and editorial workflow today; project and author visibility for the publisher; and, as the partnership develops, author development and research, cover design, composition and formatting, and downstream publishing and marketing capability. Much of the foundation is already built and working — and we would rather show you exactly which parts than round up. What is not built yet, we develop as partners, in the order your operation actually needs it.

3 · What the system does today

The production line. Every book moves through seven stations. For each one the system records what was done, who ran it, what has to be true for the book to leave, and **who closes that gate** — the author or the publisher. Not a status word; the mechanism that earned the word. The editorial stations are run by a named AI editorial team — Alex (developmental), Sam (line), Jordan (copy) — working full manuscripts, chapter by chapter, live in production today.

The book surface is the drill-down. When a title needs attention, one click gives the reason: the stations, the gates, the decisions recorded against it, the manuscript itself, the cover assets.

Decisions leave a record. A cover approval or a revision request becomes an attributed, append-only entry — who did what, when.

Everything the machine does is metered — every AI action recorded per call, per author, per manuscript, per station. Editorial work is a defined countable unit: a completed editor journey. **A pass that fails costs us and bills you nothing.**

The Lobby — the list-level answer. One screen for the whole list is **live and reads the organisation’s real data** — opened and confirmed on screen, including its designed empty state, which says “no titles yet” rather than pretending to knowledge it doesn’t have. Its attention-sort and per-imprint filters are built and take effect as titles enter the line.

4 · A phased implementation

Adoption is phased twice over, and both dials are in your hands.

Phase 1 — Editorial foundation (this proposal). The publisher workspace, your authors and projects, and the Developmental → Line → Copy editorial workflow — **running to a finished, edited manuscript and its assets, and ending, deliberately, at handoff to composition.** It starts with **two pilot titles**, one per imprint, free until the first editorial journey completes (§6).

Phase 2 — Composition & production. Turning the finished manuscript into the files your channel accepts — composition, formatting, cover typography — together with the wider production environment and author development through Wright. Explored together, with commercials agreed then, not now. §5 says plainly where this stands today.

Phase 3 — Publishing ecosystem. As High Line develops: marketing and further downstream capability, where there is a clear operational fit.

And within every phase, an authority dial: Observe → Assist → Operate, per module, switches held by your admin, not by us. At Observe the system **changes nothing about the work; it records what was done** — your editors keep working exactly as they do, and you advance a module only when you choose. The switch mechanics are in build now; until they ship, levels are set with us rather than self-served.

5 · What is live, what is in build, what we are not claiming

Live in production today: the production line, the book surface, the manuscript reader, cover artwork generation, the attributed decision record, the metering, the Lobby reading live tenancy data, and public author-side pricing at authorslab.ai/pricing. Front and back matter, listing metadata, ISBN route, pricing, channel selection and launch date are captured as structured data on the author's surface today.

In build, named honestly: the working surfaces for organisations, imprints, seats and permissions — foundations applied and their enforcement demonstrated in production, screens in progress. Per-module authority switches. Publisher-side upload. Sending a recorded note through to the author.

SECTION TO FOLLOW — staff & permissions: what a High Line admin can do on day one, in plain language. Owner: identity & billing.

Where our line ends — said exactly. Our editorial line runs to a finished, edited manuscript, its editorial record, and cover artwork ready for typography. **The manuscript compiler exists and assembles a complete book — front matter, chapters, back matter — as a single styled document; converting that into distributor-accepted files is not yet working, and what “accepted” requires depends on who receives them.** That is Phase 2's first job, and we would rather build it against the real spec than an invented one — so the question we will ask rather than answer: **what does Hachette need from you, and in what form — and who does that work for you today?**

Visibility, not forecasting — said plainly. What the system can honestly tell you today is **what has moved, what has not, and who each book is waiting on** — a stalled book with a publisher-side gate open is exactly the thing nobody at High Line can see today. What it cannot yet say is “*this book will miss March*”: there is no target date held per book in production, so where no date exists the surface says so, and a moving book is called **moving**, never “*on track*”. The next primitive on our line is the one your question asks for: a per-title target date, set by you, so that *late* has something to be measured against.

A method, not just a promise. Where a control's machinery does not exist, the control **does not appear** — **every control passes through one gate** that hides what has no substrate; we removed working-looking controls from these screens last week for exactly that reason. The same rule produced this document.

What we are explicitly not claiming. You want this embedded in the organisation — the right ambition, and *your goal, not yet our capability*. No integration with your existing systems, Hachette’s distribution machinery, contracts or royalties is claimed anywhere here. What we propose runs alongside what you have, with the record it produces as the argument for what comes next.

6 · The commercial principle

Simple, and aligned to value. Two numbers, GBP, nothing metered behind either.

Platform fee	£750 / month	The instrument: whole backlog, every seat, every imprint, at whatever authority levels you set.
Per title worked	£400, once	Fired when the system completes its first editorial station for that title. Opens the full editorial line to that book.
Pilot	Free	Two titles, one per imprint — converting only when the first editorial journey completes.

What the title fee buys: the book’s entry to the editorial line — all three editors, at your authority level, re-runs within fair use. **How much of the line you route through us is your choice; the price does not change with it.** A book loaded but never worked is never billed; a book whose launch is delayed downstream is not billed again for waiting.

Principles: no per-seat pricing; no usage meters; no overage charges; a clear relationship between payment and actual title activity. You already account per title — £400 lands as one line inside a cost structure you already run, next to human editorial at £1,600–£8,000 for a single pass on the same book. A fifteen-title year is **£15,000**; a thirty-title year **£21,000**; and if you keep your own editors throughout and use only the visibility layer, a year is **£9,000 flat** — a floor we designed, not a loophole.

7 · Proposed next step

One small decision: **agree the two pilot titles, one from Odessa, one from Antidote.**

Then, in order: we establish the High Line workspace — your organisation, your two imprints, your team’s access arriving with it; the pilot titles enter the line; the first editorial journeys run; and we review the experience together before agreeing the next phase.

You will not be implementing alone. Carl, our co-founder, will be personally available to work with your team on implementation inside High Line — setup, workflow fit, and the first journeys. And we will be asking for feedback throughout, deliberately: the point of the pilot is to shape the platform around how High Line actually runs, so the system earns its place by making your operation measurably more efficient, not by asking your operation to bend around it.

The objective is not to sell High Line software. It is to help a new publisher establish modern publishing infrastructure with as little friction as possible — and to earn each next phase by what the last one measurably did.

Working positions: UK contracting entity · GBP · invoiced net-30 — confirmed at signature. Opening terms for High Line specifically, not a rate card. Internal draft: Carl voice pass → marketing register pass → builder countersign against open screens → adversarial read → claim-by-claim verification precede any send; Paul and Carl own the send.